

Capstone Project Documentation

Cloud-Native Disaster Recovery and Infrastructure Automation using AWS

Author: Koppineni Harish Arjun

Table Of Contents

Application Stack

Part 1: Project Initialization and Prerequisite Setup

- Objective Definition
- Required Tools and Access
- Source Code Preparation

Part 2: Primary Region Infrastructure Deployment (us-east-1)

- Review CloudFormation Template
- Execute CloudFormation Stack Creation
- Verify EKS Cluster Connectivity

Part 3: Secondary Region Infrastructure Deployment (us-west-2)

- Prepare Terraform Configuration
- Execute Terraform Deployment
- Verify Secondary EKS Cluster

Part 4: CI/CD Pipeline Configuration and Automation

- Create ECR Repository
- Create Kubernetes Manifests
- Configure the buildspec.yml
- Build the AWS CodePipeline
- Static Code Analysis & Container Security (SonarQube + Trivy)
- Grant Permissions to CodeBuild Role

Part 5: Route 53 Global DNS Failover

- Create Hosted Zone
- Configure DNS Failover Records
- Health Check

Part 6: Monitoring and Alerting (CloudWatch + SNS)

- EKS Container Insights Dashboard
- CloudWatch Alarms
- SNS Setup

Part 7: Security Best Practices

- IAM and Role Isolation
- Secrets Management
- Network Security

Final Notes

Project Overview

This project implements a **highly available, fault-tolerant**, and **disaster-resilient** infrastructure for a **production-grade web application**, built on **Amazon Web Services (AWS)**. The architecture supports **multi-region failover**, fully **automated CI/CD pipelines**, and **robust monitoring and security practices** using modern DevOps tools and principles.

The infrastructure is provisioned using **Infrastructure as Code (IaC)** with **AWS CloudFormation** and **Terraform**, and integrates tightly with AWS services for resilience, scalability, and automation.

Key Features

- **Multi-Region High Availability**
 - Primary Region: us-east-1
 - Secondary Region: us-west-1
 - Disaster recovery using **Route 53 DNS Failover**
 - **End-to-End CI/CD Automation**
 - Powered by **AWS CodePipeline, CodeBuild, CodeDeploy**
 - Integrated with **ECR** for container registry
 - **Real-Time Monitoring and Alerts**
 - Metrics and logs via **Amazon CloudWatch**
 - Alerting through **Amazon SNS**
 - **Security Hardening**
 - IAM roles with least-privilege principle
 - Fine-grained **Security Groups, NACLs, and KMS encryption**
 - **DevSecOps Integrations**
 - **SonarQube** for static code analysis
 - **Trivy** for container vulnerability scanning
-

Application Stack

Tier	Technology
Frontend	Thymeleaf (HTML/CSS rendered)
Backend	Spring Boot (REST APIs + MVC)
Database	Amazon RDS MySQL (Multi-AZ)
Containerization	Docker
Orchestration	Amazon EKS (Kubernetes)
Load Balancing	Application Load Balancer (ALB)
Networking	VPC with public/private subnets, NAT Gateways, Route Tables

Part 1: Project Initialization and Prerequisite Setup

1.1. Objective Definition

The primary goal is to deploy a web application across two AWS regions (us-east-1 as Primary, us-west-1 as Secondary) to ensure high availability and robust disaster recovery capabilities. All infrastructure will be provisioned using code, and application deployments will be fully automated.

1.2. Required Tools and Access

Before beginning, ensure the following are installed, configured, and accessible:

- **AWS Account:** Administrator-level access to an AWS account.
- **AWS CLI:** Authenticated with aws configure.
- **Git & GitHub Account:** Git installed, and a personal GitHub account available.
- **Terraform:** Latest version installed.
- **kubectl:** Kubernetes command-line tool installed.
- **Docker Desktop:** For local container testing (Optional).

1.3. Source Code Preparation

1. **Fork the Repository:** Create a personal fork of the official application repository from https://github.com/arjunkoppineni/AWS_CapStone_Final_Project.git.

2. **Clone Locally:** Clone your forked repository to your local development machine.

```
git clone <your-forked-repository-url>
```

Part 2: Primary Region Infrastructure Deployment (us-east-1)

This section covers provisioning the entire infrastructure stack in the us-east-1 region using AWS CloudFormation.

2.1. Review CloudFormation Template

Navigate to the region-1-cloudformation/ directory in your cloned repository. Examine the infra.yaml file to understand the resources being created:

Networking Layer

Resource	Description
VPC (Virtual Private Cloud)	Custom VPC created for the application in each region with a /16 CIDR block (e.g., 10.0.0.0/16), providing private IP addressing for all AWS resources.
Subnets (6 total)	Split across 2 Availability Zones for high availability: - 2 Public Subnets (1 per AZ) for ALB and NAT Gateways - 4 Private Subnets (2 per AZ) for EKS worker nodes and RDS databases
Internet Gateway (IGW)	Attached to the VPC to allow internet access from the public subnets (e.g., for ALB).
NAT Gateways (2)	Placed in each AZ's public subnet, allowing EKS worker nodes and private RDS instances to securely reach the internet (for patching, pulling container images, etc.).
Route Tables	Custom route tables for public and private subnets: - Public subnet route table forwards traffic to the IGW. - Private subnet route tables forward outbound traffic through NAT gateways.

Compute Layer: Amazon EKS

Resource	Description
Amazon EKS Cluster	Fully managed Kubernetes control plane (1 per region) that orchestrates pods and deployments for the Spring Boot application.
EKS Node Group (Managed)	Auto-scaled group of EC2 instances (typically t3.medium or m5.large) managed by EKS for compute workloads. It spans multiple subnets for HA.
Security Groups for EKS	Controls inbound/outbound traffic: - Allow only ALB-to-node communication (e.g., ports 80/443) - Nodes can access RDS (port 3306) - EKS worker nodes have open internal cluster communication within VPC

IAM (Identity & Access Management)

Role	Usage & Permissions
EKS Cluster Role	Grants permissions for EKS control plane to interact with other AWS services (e.g., ELB, Auto Scaling, etc.). - AmazonEKSClusterPolicy, AmazonEKSServicePolicy
EKS Node IAM Role	Attached to the EC2 worker nodes to allow access to Amazon ECR and other AWS resources. - AmazonEC2ContainerRegistryReadOnly, AmazonEKSWorkerNodePolicy, AmazonEKS_CNI_Policy, CloudWatchAgentServerPolicy
CodeBuild Role	Executes pipeline build jobs. Has permissions to: - Build/push Docker images - Pull Secrets - Update EKS (via kubectl) - Needs inline policy for eks:DescribeCluster, eks:UpdateConfigMap - Needs AmazonEC2ContainerRegistryPowerUser
CodePipeline Role	Coordinates actions between GitHub (source), CodeBuild (build), and EKS (deploy). Includes permissions for: - codecommit:*, codebuild:*, eks:*, s3:*
CloudFormation Execution Role	Created during stack deployment. Allows CFN to create/modify AWS resources.

2.2. Execute CloudFormation Stack Creation

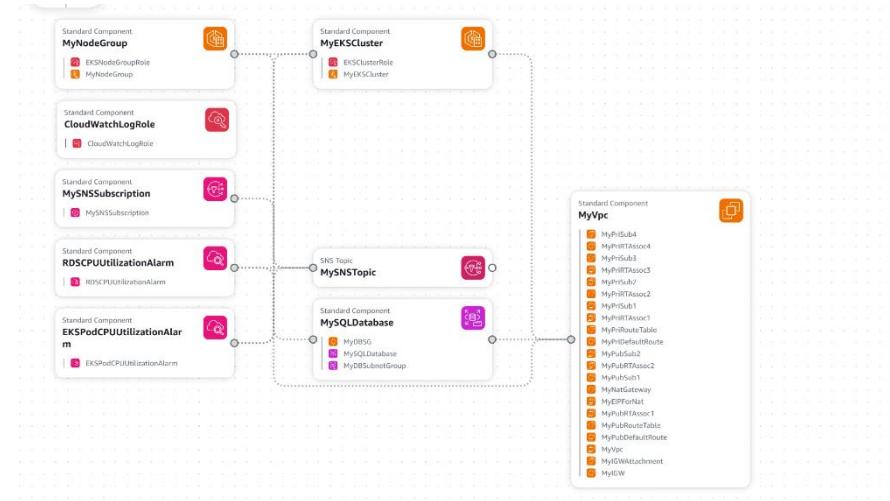
1. **Navigate to AWS Console:** Log in and go to the CloudFormation service in the **N. Virginia (us-east-1)** region.
2. **Initiate Stack Creation:** Click **Create stack > With new resources (standard)**.
3. **Upload Template:** Select **Upload a template file** and choose the infra.yaml file.
4. **Specify Parameters:**
 - o **Stack name:** Provide a unique name (e.g., prod-east-stack).
 - o **Parameters:** Carefully fill in the EKS cluster name, database name, and master credentials. **Securely store the database password.**
5. **Acknowledge and Create:** Proceed through the options, acknowledge the IAM resource creation at the final review step, and click **Create stack**. (**Note:** This process takes approximately 20 minutes.)

2.3. Verify EKS Cluster Connectivity

1. Once the stack status is CREATE_COMPLETE, configure kubectl to communicate with the new cluster.
2. `aws eks update-kubeconfig --region us-east-1 --name <your-eks-cluster-name>`
3. Confirm that worker nodes are connected and ready.
4. `kubectl get nodes`

The output should list two nodes with a Ready status.

This are the resources that are created from the template



VPC ID: vpc-02237f59bf0414123

Details

- State:** Available
- Tenancy:** default
- Default VPC:** No
- IPv4 CIDR:** 10.0.0.0/16
- IPv6 CIDR:** -
- Network Address Usage metrics:** Disabled
- Block Public Access:** Off
- DHCP option set:** dopt-057aba29a0c899913
- Main route table:** rtb-0a3b6fc0017834804
- IPV6 pool:** -
- Route 53 Resolver DNS Firewall rule groups:** -
- DNS hostnames:** Enabled
- Owner ID:** 216989142685

Resource map

- VPC:** Capstone-CF-Infra-My-VPC
- Subnets (6):** Capstone-CF-Infra-Public-Subnet..., Capstone-CF-Infra-Private-Subnet..., Capstone-CF-Infra-Private-Subnet...
- Route tables (3):** rtb-0a3b6fc0017834804, Capstone-CF-Infra-Public-RT, Capstone-CF-Infra-Private-RT
- Network connections (2):** Capstone-CF-Infra-Internet-Gateway, Capstone-CF-Infra-NAT-Gateway

MyEKSCluster

Cluster info

- Status: Active
- Kubernetes version: 1.32
- Support period: Standard support until March 21, 2026
- Provider: EKS

Cluster health

- Upgrade insights: 4
- Node health issues: 0

Overview

Details

- API server endpoint: https://2A14235BFDB70CA576405CF4410FB759.gr7.us-east-1.eks.amazonaws.com
- OpenID Connect provider URL: https://oidc.eks.us-east-1.amazonaws.com/id/2A14235BFDB70CA576405CF4410FB759
- Created: a day ago

Part 3: Secondary Region Infrastructure Deployment (us-west-1)

This section details the provisioning of the failover infrastructure in us-west-1 using Terraform.

3.1. Prepare Terraform Configuration

Navigate to the terraform/ directory. The configuration files here mirror the CloudFormation setup using Terraform

3.2. Execute Terraform Deployment

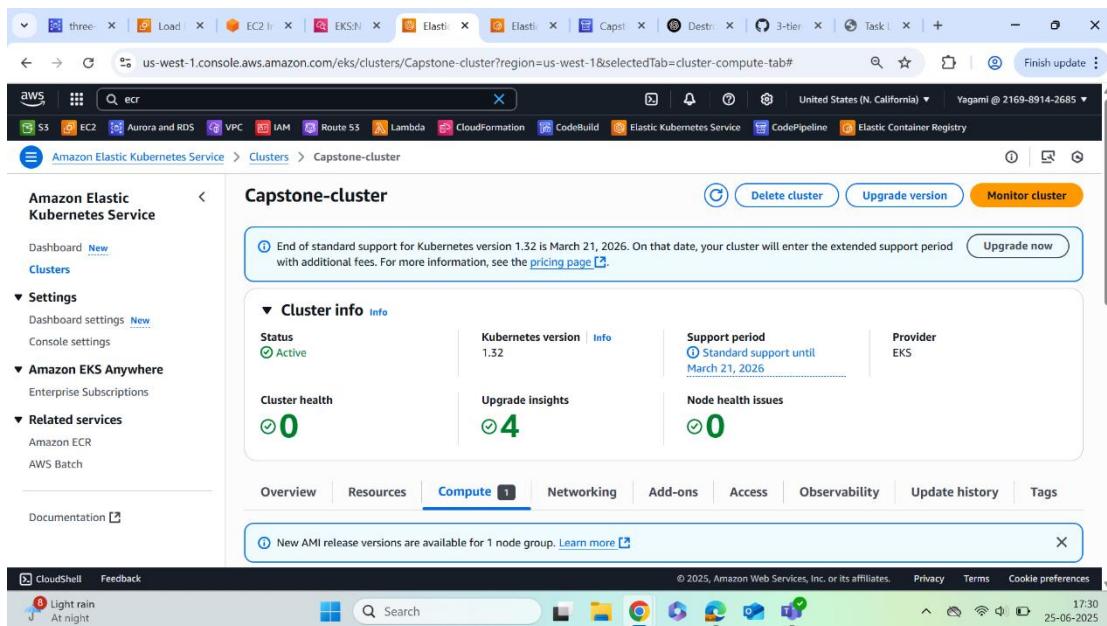
- Initialize Terraform:** Open a terminal in the region-2-terraform directory and run:
- terraform init**
- Plan Deployment:** (Optional but Recommended) Preview the resources created.

4. terraform plan
5. **Apply Configuration:** Provision the infrastructure. Type yes when prompted to
6. terraform apply (**Note:** This process also takes approximately 20 minutes.)

3.3. Verify Secondary EKS Cluster

1. Configure kubectl for the new us-west-1 cluster.
2. aws eks update-kubeconfig --region us-west-1 --name <your-region-2-CN>
3. kubectl get nodes

(Tip: Use `kubectl config get-contexts` to see and `kubectl config use-context <name>` to switch between your primary and secondary clusters.)



Part 4: CI/CD Pipeline Configuration and Automation

This section details the setup of an automated pipeline to build and deploy the application to the primary EKS cluster.

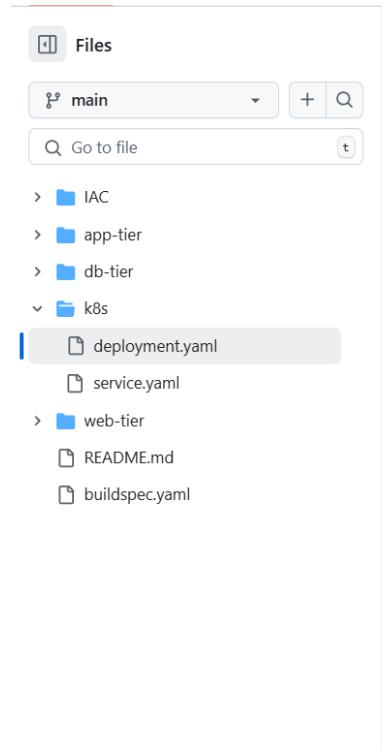
4.1. Create ECR Repository

1. In the AWS Console, navigate to **Elastic Container Registry (ECR)** in the us-east-1 region.
2. Create a new **private** repository. Name it my-app or similar.
3. Note the **repository URI** for later use.

4.2. Create Kubernetes Manifests

In the root of your repository, create a k8s/ directory and add two files:

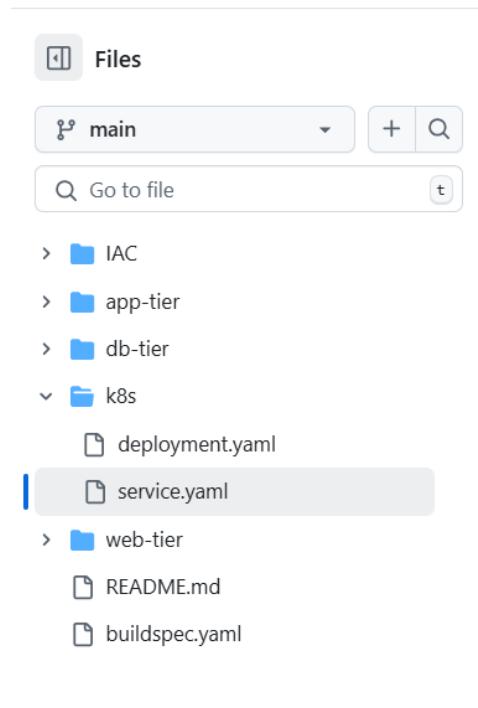
- **k8s/deployment.yaml:**



The screenshot shows the AWS CodeCommit interface. On the left, there's a sidebar with 'Files' and a search bar. Below it is a tree view of the repository structure: IAC, app-tier, db-tier, k8s (which contains deployment.yaml, service.yaml, README.md, and buildspec.yaml). The 'deployment.yaml' file is selected. On the right, the main panel shows the file content with a 'Code' tab selected. The code is a Kubernetes Deployment manifest.

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: todo-app
5    labels:
6      app: todo-app
7  spec:
8    replicas: 2
9    selector:
10      matchLabels:
11        app: todo-app
12    template:
13      metadata:
14        labels:
15          app: todo-app
16      spec:
17        containers:
18          - name: todo-app
19            image: image_url
20            imagePullPolicy: Always # <-- This ensures the image is pulled every time
21        ports:
22          - containerPort: 80
23        env:
```

k8s/service.yaml:



The screenshot shows the AWS CodeCommit interface. On the left, there's a sidebar with 'Files' and a search bar. Below it is a tree view of the repository structure: IAC, app-tier, db-tier, k8s (which contains deployment.yaml, service.yaml, README.md, and buildspec.yaml). The 'service.yaml' file is selected. On the right, the main panel shows the file content with a 'Code' tab selected. The code is a Kubernetes Service manifest.

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: todo-app-service
5  spec:
6    selector:
7      app: todo-app
8    ports:
9      - protocol: TCP
10        port: 80
11        targetPort: 80
12    type: LoadBalancer
```

4.3. Configure the buildspec.yml

Update the buildspec.yml file in your repository to contain the complete build, push, and deploy logic. Ensure the placeholder DOCKER_IMAGE_PLACEHOLDER is used.

4.4. Build the AWS CodePipeline

1. In the AWS Console, navigate to **CodePipeline** and start the creation wizard.
2. **Stage 1: Source:** Connect to your forked GitHub repository (main branch).
3. **Stage 2: Build:**
 - o Select **AWS CodeBuild** as the provider.
 - o Click **Create project** and configure it:
 - **Environment:** Use a managed **Amazon Linux 2** image.
 - **Privileged:** **Enable** this flag for building Docker images.
 - **Service Role:** Create a new service role. **You must edit this role later.**
 - **Buildspec:** Use the buildspec.yml from your source code.
 - o Add the **ECR_REPOSITORY_URI** and **EKS_CLUSTER_NAME** as environment variables in the CodeBuild project configuration.

Static Code Analysis & Container Security

To ensure high code quality and secure container images before deploying to production, the CI/CD pipeline integrates:

SonarQube – Static Code Quality Checks

- **Purpose:** Analyzes codebase for code smells, bugs, vulnerabilities, and maintainability.
- **Integration:**
 - o SonarQube is triggered during the build stage of CodeBuild.
 - o SonarQube scanner is configured via environment variables: SONAR_TOKEN, SONAR_HOST_URL.
 - o Fails the build if quality gates are not met.

- **SonarQube Scan** : Make sure SonarQube server is publicly reachable (or hosted in VPC) and credentials are securely stored in AWS Secrets Manager or CodeBuild environment variables.

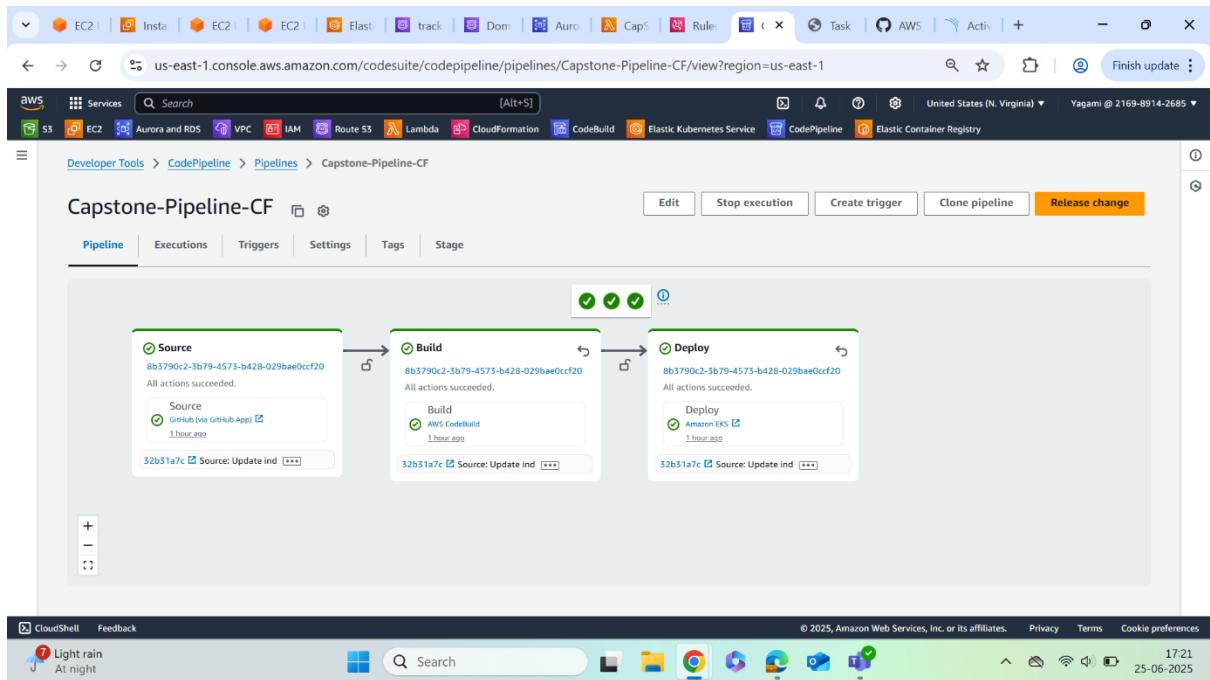
Trivy – Container Vulnerability Scanning

- **Purpose:** Scans container images for vulnerabilities (CVEs), misconfigurations, and secrets before pushing to ECR.
 - **Integration:**
 - Used in CodeBuild phase after image build, before pushing to ECR.
 - Fails the build if critical vulnerabilities are found.
 - **Trivy Scan**
4. **Stage 3: Deploy: Deploy to EKS**, It Deploys the new Docker image to the EKS cluster by applying updated Kubernetes manifests.
 5. **Create** the pipeline.

4.5. Grant Permissions to CodeBuild Role

1. Navigate to **IAM > Roles**.
2. Find the service role created for your CodeBuild project.
3. Attach the necessary permissions:
 - AmazonEC2ContainerRegistryPowerUser (to push to ECR).
 - An inline policy allowing eks:DescribeCluster.
4. **Crucially, authorize the role in EKS:** Add the CodeBuild role's ARN to the aws-auth ConfigMap in your us-east-1 EKS cluster to grant it kubectl permissions.

After commit to GitHub → Pipeline triggers → CodeBuild builds, pushes, deploys.



These steps should be replicated in Region 2 (us-west-2) to ensure consistent infrastructure and deployment configuration

Part 5: Route 53 Global DNS Failover

5.1 Create Hosted Zone

1. Go to Route 53 → Hosted zones → Create hosted zone
2. Enter domain (use registered one or subdomain for test)

5.2 Configure DNS Failover Records

- **Record A (us-east-1)**
 - Alias to ALB DNS in Primary Region
 - Failover: Primary
 - Attach Health Check for ALB
- **Record A (us-west-1)**
 - Alias to ALB DNS in Secondary Region
 - Failover: Secondary

5.3 Health Check

- Target: ALB public endpoint
- Protocol: HTTP, Port 80 or 443
- If health check fails → Route 53 redirects to Secondary ALB

The screenshot shows the AWS Route 53 Health checks console. On the left, a sidebar navigation includes: Dashboard, Hosted zones, Health checks (selected), Profiles (New), IP-based routing, Traffic flow, Domains, and Resolver. The main content area displays a table titled "Health checks (2) Info". The table lists two endpoints:

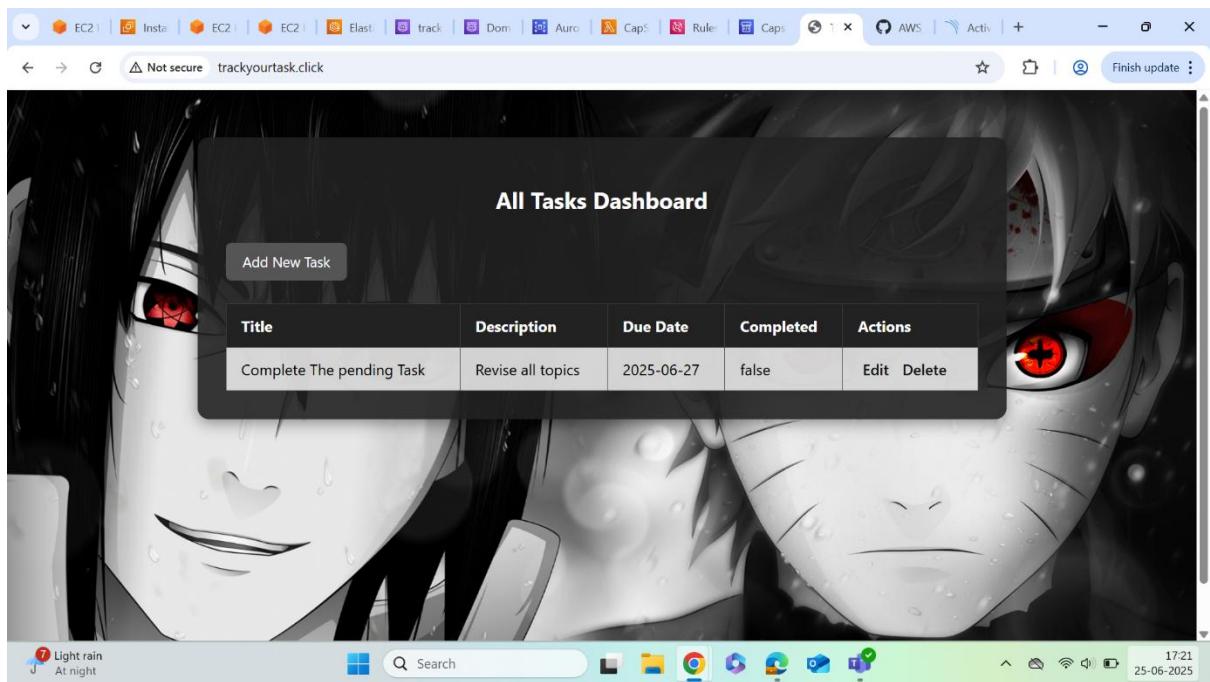
ID	Name	Details	Status in last 24 hours	Current s...	Alarm	Actions
5780511e-0ea8...	us-east-1	http://ac2b620f...	<div style="width: 100%; height: 10px; background-color: green;"></div>	Healthy	None, ...	⋮
f6c24270-554f...	us-west-1	http://a512785...	<div style="width: 100%; height: 10px; background-color: green;"></div>	Healthy	None, ...	⋮

A message at the top states: "Introducing the new Route 53 health checks console experience. We've redesigned the health checks console to make it easier to use. The changes include a new layout for faster access to information. Learn more about the changes and let us know what you think. Or you can use the old console." A "Create health check" button is located in the top right of the main area.

The screenshot shows the AWS Route 53 Hosted zones console. The sidebar navigation is identical to the previous screen. The main content area displays a table titled "Hosted zone details" for the zone "trackyourtask.click". The table has three tabs: "Records (4)", "DNSSEC signing", and "Hosted zone tags (0)". The "Records (4) Info" section indicates "Automatic mode is the current search behavior optimized for best filter results. To change modes go to settings." The table lists four records:

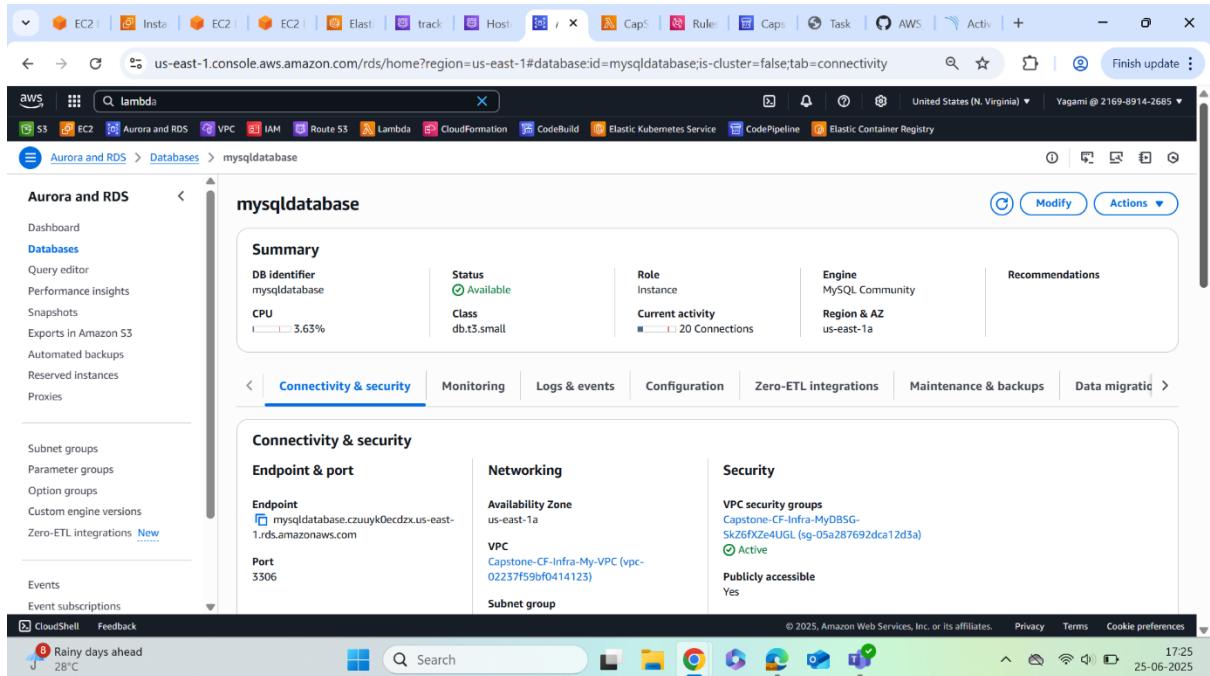
Type	Routing p...	Alias	Value/Route traffic to	TTL (s...)	Health ...	Evalu...	Record
A	Failover	Primary	dualstack.ac2b620f67ee44c...	-	f6c24270...	Yes	east...
A	Failover	Secondary	dualstack.a512785548df94c...	-	f6c24270...	Yes	us-w...
NS	Simple	-	ns-1264.awsdns-30.org.	172800	-	-	ns-694.awsdns-22.net.
SOA	Simple	-	ns-370.awsdns-46.com.				ns-1773.awsdns-29.co.uk.
			ns-1264.awsdns-30.org. aws...	900	-	-	

A "Create record" button is located in the top right of the table area. The bottom of the screen shows the AWS navigation bar and a status bar indicating "CloudShell Feedback", "28°C Mostly cloudy", and the date "25-06-2025".



To simulate disaster recovery, you can intentionally fail the primary region's health check or temporarily disable its ALB. Route 53 will automatically shift traffic to the secondary region ALB.”

After accessing the application, You can see that the data stored in the RDS.



```

Tables in todo
+-----+
| to do |
+-----+
1 row in set (0.00 sec)

mysql> select * from to_do;
Empty set (0.00 sec)

mysql> select * from to do;
+----+-----+-----+-----+-----+
| id | completed | description | due_date | title |
+----+-----+-----+-----+-----+
| 4 | 0x00     | REVise all AWS topics | 2025-06-27 | Complete Certification |
+----+-----+-----+-----+-----+
1 row in set (0.01 sec)

mysql>
mysql> select * from to do;
+----+-----+-----+-----+-----+
| id | completed | description | due_date | title |
+----+-----+-----+-----+-----+
| 5 | 0x00     | Revise all topics | 2025-06-27 | complete The pending Task |
+----+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql>

```

i-098aa935e7178cc8e (cap)
Public IPs: 54.84.224.28 Private IPs: 10.0.55.146

Part 6: Monitoring and Alerting (CloudWatch + SNS)

6.0 EKS Container Insights Dashboard

AWS CloudWatch Container Insights can be enabled for EKS cluster (us-west-1) to monitor performance and operational health of pods, containers, and nodes.

Container Insights Metrics Available:

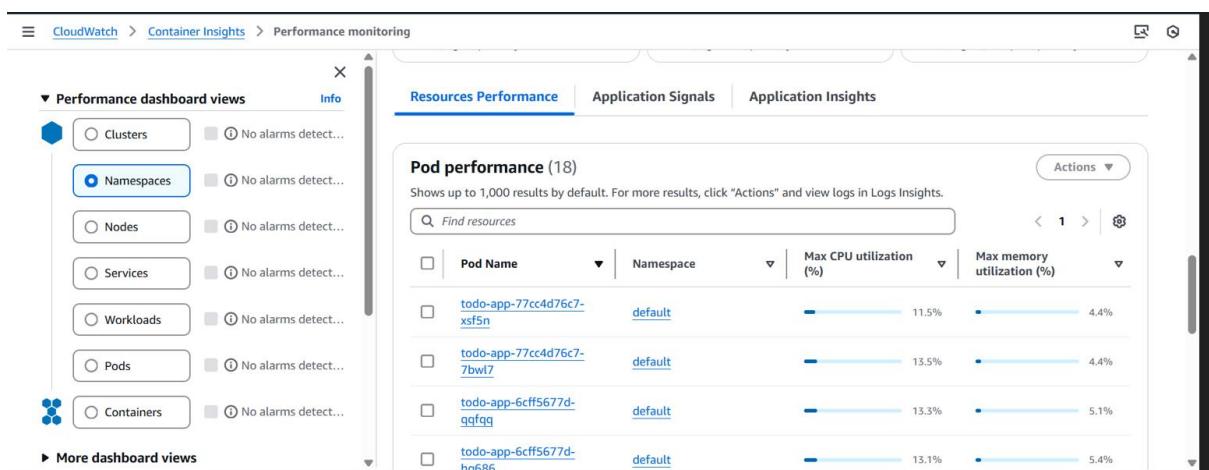
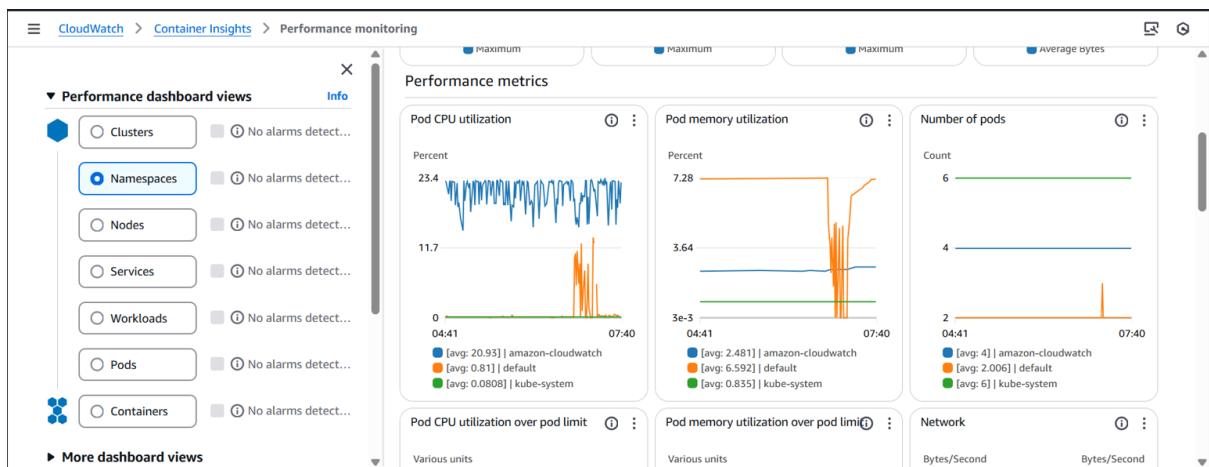
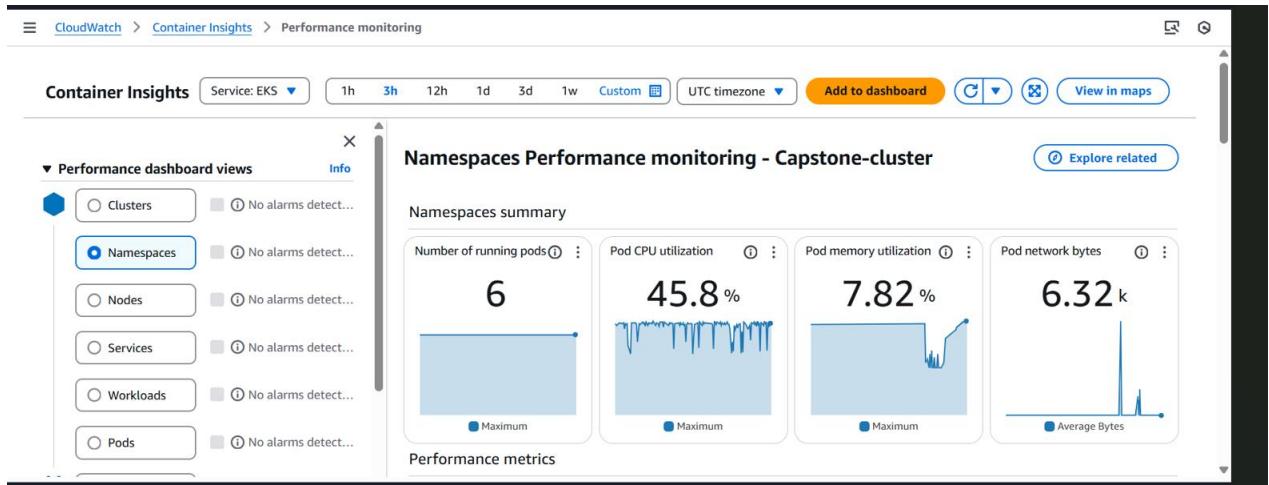
- pod_cpu_utilization
- container_memory_utilization
- node_cpu_utilization
- network_rx_bytes
- network_tx_bytes

Default CloudWatch Dashboards:

After enabling, navigate to: **CloudWatch → Dashboards → Container Insights → Performance Monitoring**

You will find:

- Node/Pod maps with CPU/Memory visualizations
- Filterable graphs for namespace, cluster, deployment, etc.
- Top N resource consumers



6.2 CloudWatch Alarms

Resource Metric	Threshold Action
RDS	CPUUtilization > 70% Notify via SNS

6.2 SNS Setup

1. Go to SNS → Create Topic (e.g., infra-alerts)
2. Add email subscription
3. Attach SNS topic to CloudWatch alarms

Part 7: Security Best Practices

7.1 IAM and Role Isolation

- Principle of Least Privilege (PoLP)
- Separate roles for CodeBuild, EKS, CloudFormation

7.2 Secrets Management

- Store DB passwords, JWT secrets in **AWS Secrets Manager**
- Alternatively, use **SSM Parameter Store (SecureString)**

7.3 Network Security

- RDS in private subnets
- Use Security Groups + NACLs to restrict access
- No open ports to public except ALB (443, 80)

Final Notes

- The application runs in **two AWS regions** (us-east-1 primary, us-west-2 secondary) for **high availability and disaster recovery**.
- **CloudFormation** and **Terraform** are used for fully automated infrastructure setup.
- **CI/CD** is implemented using **CodePipeline**, with **SonarQube** for code analysis and **Trivy** for image scanning.
- **Route 53** provides DNS-based **failover** between regions.
- **CloudWatch** and **SNS** handle monitoring and alerts.
- **Security** is enforced using IAM roles, Secrets Manager, and private networking.
- The setup is **automated, secure, and scalable** for production use.